



High Level Design (HLD) - Mini SOC Platform



Overview

This document describes the high-level design for the Mini SOC (Security Operations Center) platform using Wazuh SIEM (Security Information and Event Management), docker swarm (container orchestration), and DevOps practices.

Scope

- Deployment of Wazuh stack on Docker Swarm
- high availability (HA), disaster recovery (DR), fault tolerance, and scalability

Wazuh SIEM Overview

What is Wazuh?


- Wazuh is an **open-source security platform** that provides **log analysis**, **intrusion detection**, **vulnerability detection**, and **compliance monitoring**.
- It acts as a **Security Information and Event Management (SIEM)** system, collecting and analyzing data from endpoints, servers, and network devices

Key Components

- **Wazuh Manager:** Central component that analyzes **security events**, generates **alerts**, and manages **agents**.
- **Wazuh Agents:** Installed on endpoints or servers to collect logs, monitor files, detect anomalies, and send data to the manager.
- **Elasticsearch / Indexer:** Stores and indexes the collected data for fast search and analytics.
- **Kibana:** Provides dashboards and visualizations for monitoring and reporting.
- **Filebeat / Log Shippers:** Forward logs from agents to Elasticsearch in real time.

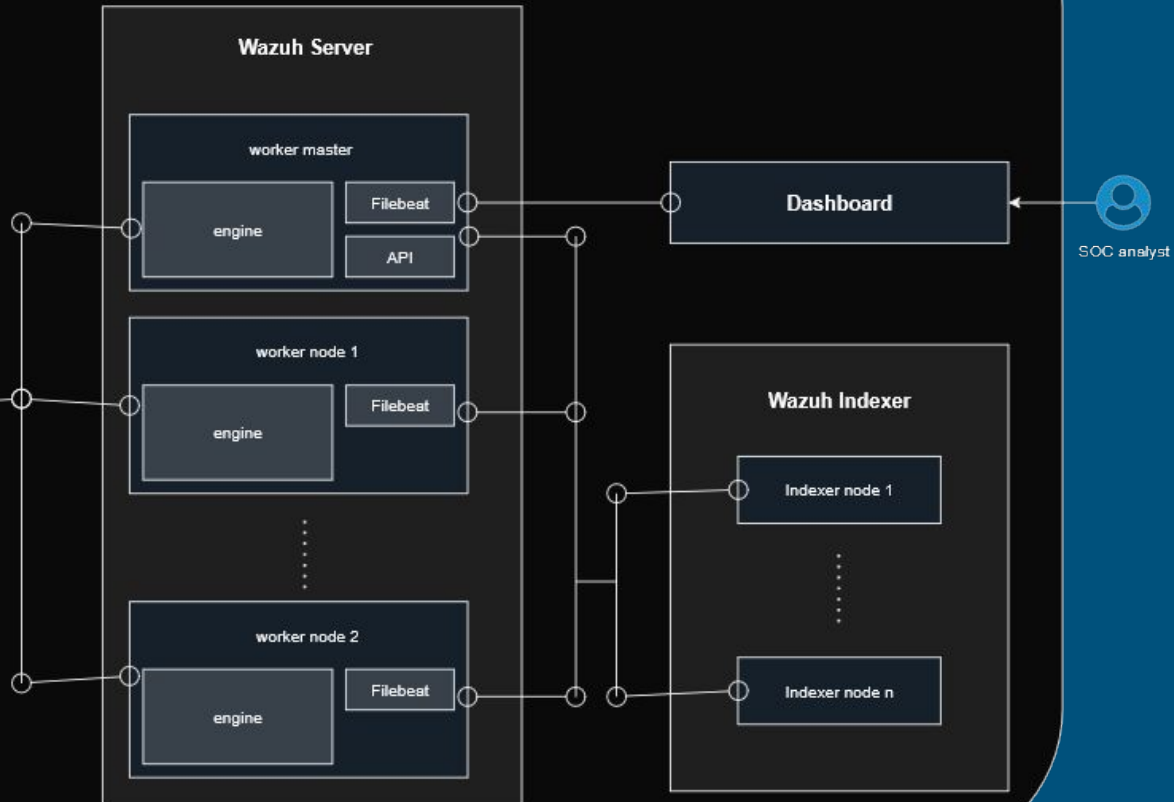
Wazuh stack Architecture

Endpoints




Load Balancer

Wazuh Central Components



Container Orchestration with Docker Swarm



Docker Swarm Overview

- **Container orchestration tool built into Docker**

Enables you to manage multiple containers across different hosts as a single cluster.

Unlike standalone Docker, Swarm automates deployment, scaling, and container lifecycle management.

- **Manages clusters of Docker nodes (machines) as a single virtual system**

Nodes (physical or virtual machines) are grouped under one “swarm,” appearing as one logical Docker host.

This allows distributed workloads to be managed from a single control point.

- **Provides high availability, scalability, and load balancing**

Swarm automatically replicates services across nodes to prevent downtime.

Traffic is distributed evenly between containers using built-in load balancing.

You can easily scale services up or down with one command (`docker service scale`).

- **Secure and integrated by design**

Uses mutual TLS for secure node communication.

Works natively with existing Docker CLI and Compose files, requiring minimal new tooling.



How It Works

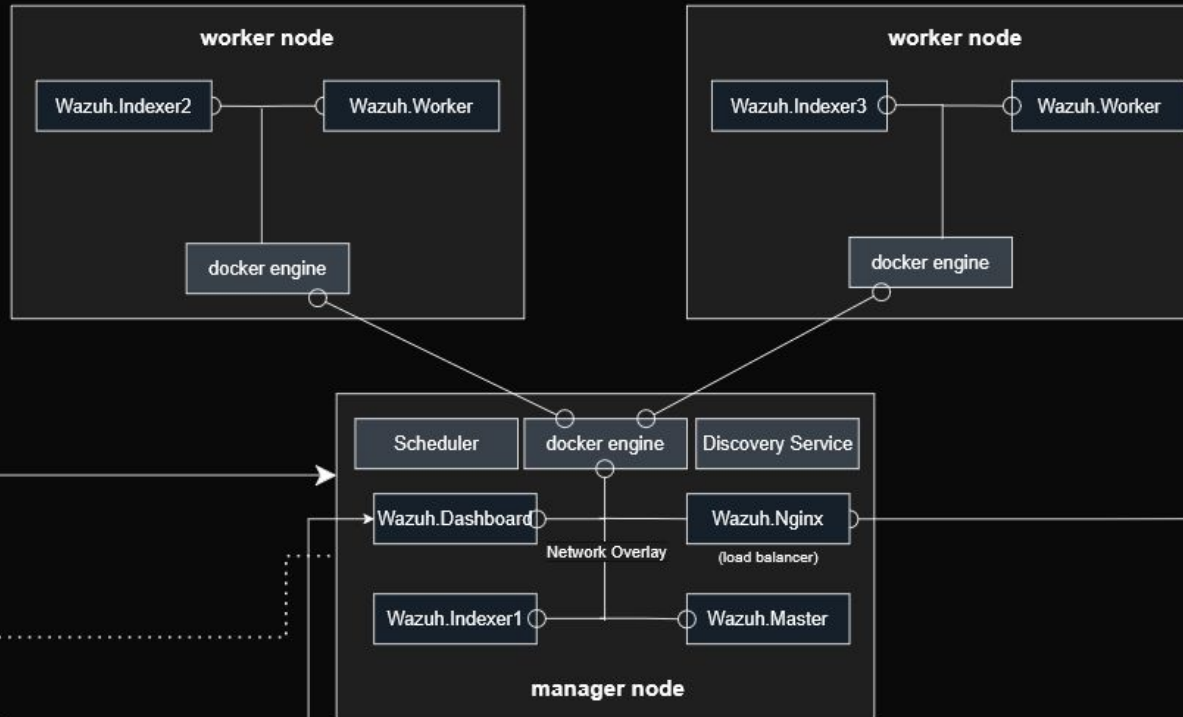
- **Manager Nodes:** Control the cluster, schedule services, handle orchestration.
- **Worker Nodes:** Run the containers (services) assigned by the managers.
- **Overlay Networks:** Enable container communication across multiple hosts.
- **Service Replication:** Automatically restarts failed containers.

Why Use Docker Swarm in This Project

- Simplifies multi-container deployment (e.g., Wazuh stack).
- Ensures high availability for critical services.
- Seamlessly scales horizontally when load increases.
- Built-in integration with Docker CLI (no external setup like Kubernetes).

Wazuh Stack Architecture on Docker Swarm (3-Node Cluster)

Docker Swarm Cluster



High Availability (HA), Disaster Recovery (DR), Fault Tolerance & Scalability

Ensuring continuous operation and resilience of the Wazuh Stack on Docker Swarm.

Introduction

Why Reliability Matters

- Security platforms must operate 24 / 7 / 365.
- Outages or data loss reduce visibility and response capability.
- This section covers the design strategies that make the environment resilient, recoverable, and scalable.

High Availability (HA)

Definition:

Ability of the system to remain operational despite node or service failures.

Implementation:

- Multiple Swarm nodes (1 manager + 2 workers) for cluster redundancy.
- Wazuh Indexer replicated across nodes for continuous log indexing.
- Wazuh Nginx acts as a load balancer for dashboard/API traffic.
- Overlay networks ensure service communication even if one node fails.

Disaster Recovery (DR)

Definition:

Process for restoring services and data after a catastrophic event.

Implementation:

- Automated backups of Indexer data and configurations.
- Off-site or remote storage of backups (e.g., S3 / NAS / remote volume).
- Recovery scripts or CI/CD jobs to redeploy Swarm stacks quickly.
- Option for secondary Swarm cluster or DR site replication.

Fault Tolerance

Definition:

Capability to continue operation when individual components fail.

Implementation:

- Swarm service replication keeps containers running if one crashes.
- Health-check & restart policies automatically recover failing services.
- Decoupled services (Manager, Indexer, Dashboard) prevent cascade failure.
- Persistent volumes protect stored data during container restarts.

Scalability

Definition:

Ability to increase capacity or performance as demand grows.

Implementation:

- **Horizontal scaling:** add worker nodes to the Swarm cluster.
- **Vertical scaling:** allocate more CPU/RAM to existing nodes.
- **Dynamic scaling via CI/CD:** pipelines deploy additional replicas automatically.
- **Load balancing** distributes traffic among Indexer and Worker nodes.